Data Flash Back

Human error has always been the most difficult type of error from which to recover. This is because as far as the database is concerned, human error is not an error at all. The “error” is just another committed transaction, and the rules of a relational database (the “D” for “Durable” of the ACID test) do not allow Oracle to back out committed transactions. Depending on the nature of the error, the different Flashback technologies may help you to recover, while minimizing downtime and loss of data.

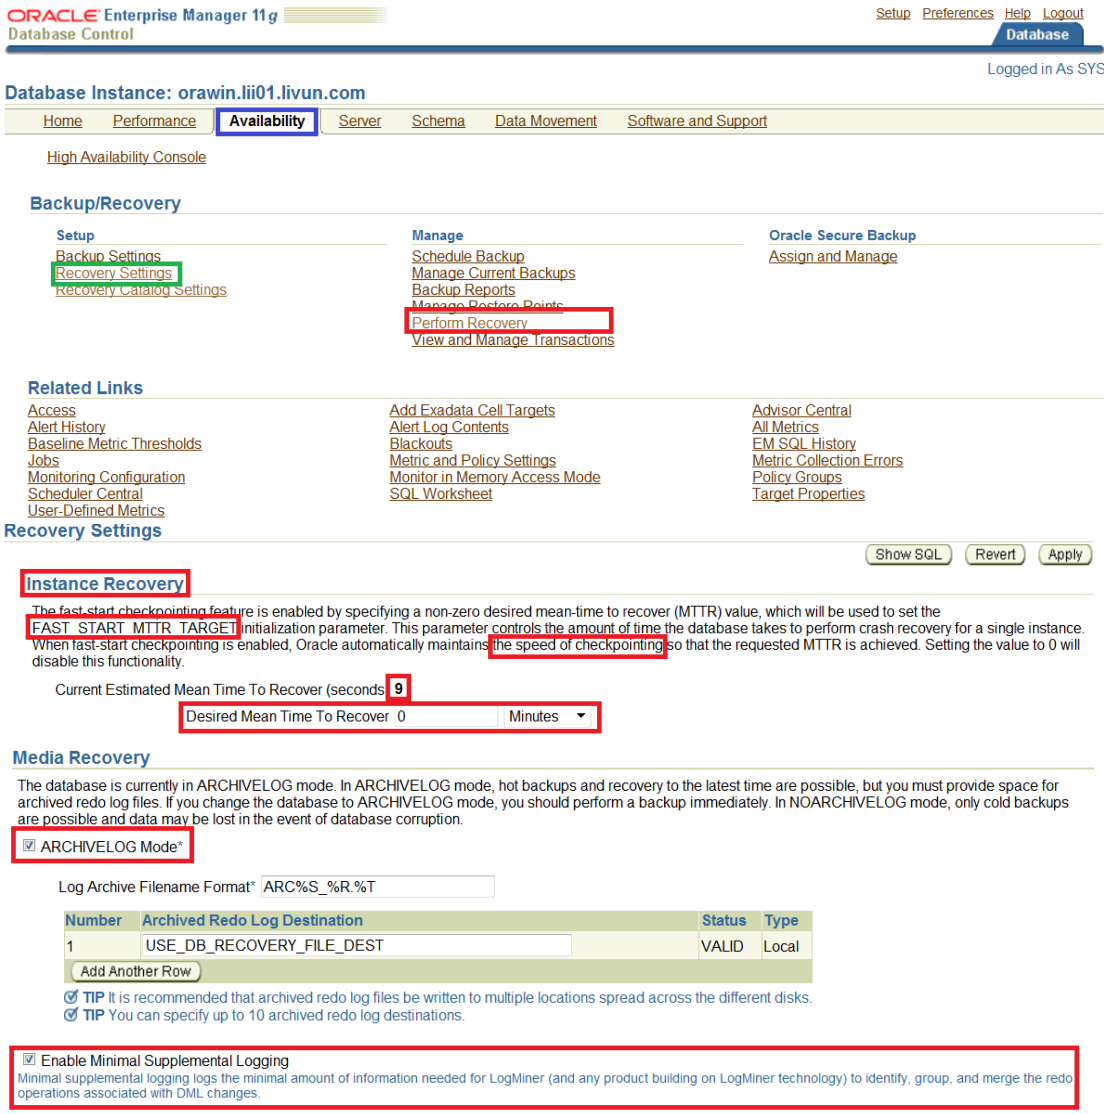

The typical reason for using any type of flashback technology is to correct mistakes.There are four distinct flashback technologies available, each implemented with a different underlying architecture. Each technology has different capabilities and limitations, but there is overlapping functionality between them. The typical reason for using any type of flashback technology is to correct mistakes—it is vital to understand what type of flashback technology is appropriate for correcting different types of errors.

Flashback Database: Flashback Database is, by analogy, like pressing a rewind button on the database. The current database is taken as the starting point, and it is taken back in time, change by change, reversing all work done sequentially. The end result is as if you had done an incomplete recovery: all

work subsequent to the flashback point is lost, and indeed the database must be opened with RESETLOGS.

Flashback Database is an alternative to incomplete recovery following user error—perhaps a more flexible and much faster alternative. Flashback Database is in fact a combination of several processes and data structures. First, you must allocate some memory in the SGA (which will be automatic - you cannot control how large the buffer is) and some space on disk to store the flashback data, and start the RVWR process to enable flashback logging. When doing a flashback, Oracle will use the flashback logs to take the database back in time to before the time you want, and then apply redo logs (using whatever archive redo log files and online redo log files are necessary) in the usual fashion for incomplete recovery to bring the datafiles forward to the exact time you want. Then the database can be opened with a new incarnation, in the same manner as following a normal incomplete recovery.

Flashback Database is restricted by the size of the `flashback logs`. Unlike the redo log, flashback logging is not a log of changes—it is a log of complete block images.

```
SQL> alter system set db_recovery_file_dest='/flash_recovery_area';
SQL> alter system set db_recovery_file_dest_size=8G;
SQL> alter system set db_flashback_retention_target=240;
SQL> shutdown immediate;
SQL> startup mount;
SQL> alter database open;
```

The flash recovery area is a default location for all recovery-related files: either permanent or transient. The permanent files are:

- Controlfile multiplexed copies
- Online redo log file multiplexed copies

The transient occupants of the flash recovery area are

- Archive redo log files
- RMAN backups
- RMAN copies
- Controlfile auto-backups
- Flashback log files

The flash recovery area is controlled by two instance parameters:

- DB_RECOVERY_FILE_DEST
- DB_RECOVERY_FILE_DEST_SIZE

```
SQL> show parameters recovery
NAME                                 TYPE        VALUE
------------------------------------ ----------- ------------------------------
db_recovery_file_dest            string      C:\app\oracle\recovery_area
db_recovery_file_dest_size      big integer 3852M
```

The first of these parameters nominates the location. The second parameter limits the maximum amount of space in the destination that the database will take up, Oracle Corporation’s advice is that the flash recovery area should be greater than the sum of the sizes of the database, all incremental backups since the last full backup, and all archive redo log files that have not been backed up to tape.

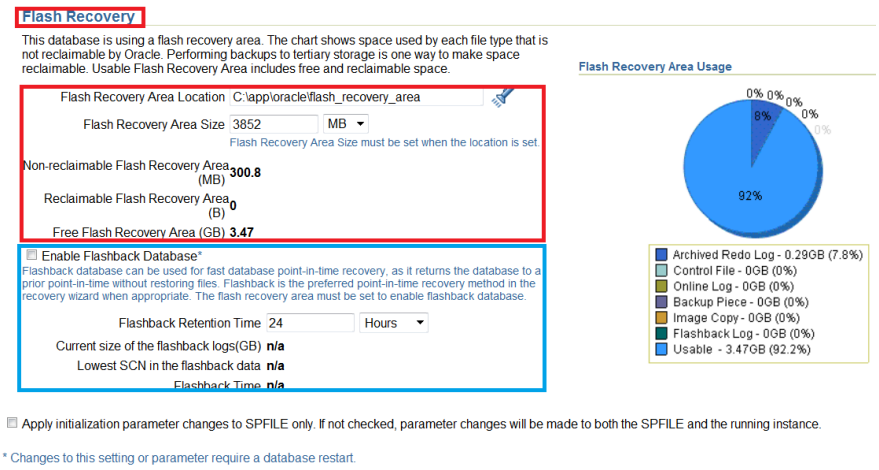

Flashback Query, Transaction, and Table: three flashback techniques are based on the use of undo segments. The first flashback capability was initially introduced with release 9i of the database and has been substantially enhanced subsequently. Flashback Query (the release 9i feature) lets you query the database as it was at some time in the past, either for one select statement or by taking your session temporarily back in time so that all its queries will be against a previous version of the database. Flashback Transaction automates the repair process. Once you have used Flashback Query to identify which transaction it was that caused the problem, Oracle can construct SQL statements that will reverse the changes. Having determined that inappropriate work has been committed against one table, you can instruct Oracle to reverse all changes made to that table since a particular point in time, while leaving all other tables current.

Throughout any Flashback Query, Transaction, or Table operation, the database remains open and all objects (including those involved in the flashback) are available for use. Transactional integrity and constraints are always enforced, which means that the flashback operation might fail. For example, if a flashback of a transaction requires an insert into a primary key column, that value must not be in use. Flashing back one table may not be possible if it has foreign key constraints—you will have to flash back all the related tables in one operation.

The flashback query is capable and limited by the undo retention:

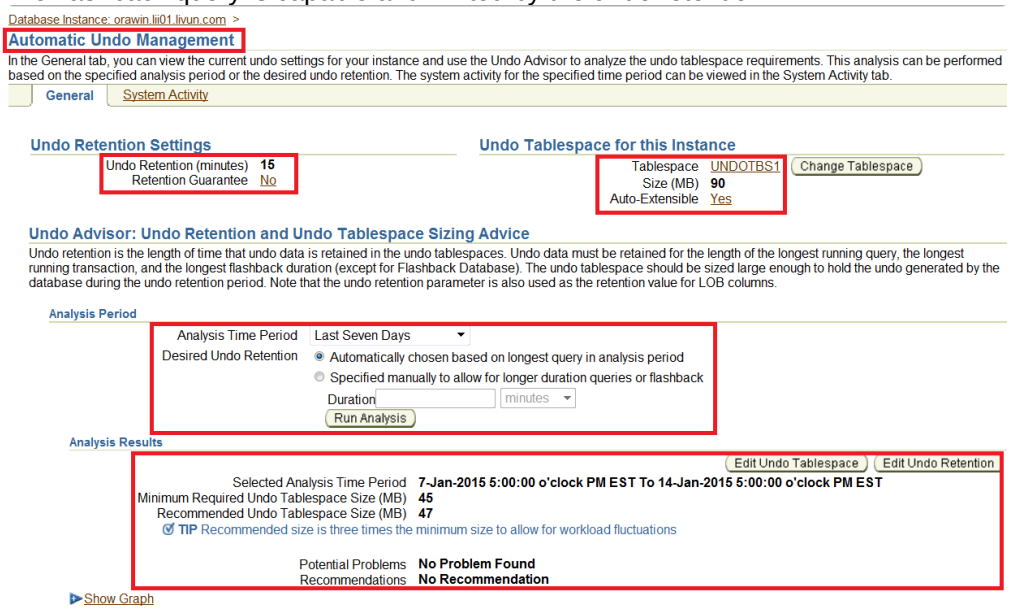

To guarantee that a flashback query will always succeed for a given period, set the `RETENTION GUARANTEE` attribute for the undo tablespace, in conjunction with the UNDO_RETENTION instance parameter. This will ensure that you can always flash back the number of seconds specified—but the price you will pay is that if your undo tablespace is not sized adequately for the transaction workload, then the database may hang while performing DML, you must monitor the `V$UNDOSTAT` view to calculate the necessary size.

Flashback Drop, It is now possible to “un-drop” a table. This is implemented by mapping the DROP command onto a RENAME command. Rather than dropping the table, it is renamed to a system-generated name, and only actually dropped later, when its storage space is needed for a live object. If necessary, and if its storage space has not been reused, the object can be renamed back to its original name and thus restored. Without this capability, the only way to get a table back after a drop would be to do an incomplete recovery to the point in time just before the table was dropped. This was usually timeconsuming, and it meant the loss of all work done subsequently. The new Flashback Database capability achieves the same result as incomplete recovery and should be much faster, but work done on other tables following the drop is lost and the database will be unavailable until the operation is completed.

```
SQL> drop table <table_name> purge;
SQL> flashback table <table_name> to before drop rename to <new_name>;
```

Flashback Drop lets you reinstate the table as it was at the time that it was dropped, with no loss of data whatsoever; the database remains current. This does not require any use of backups, and neither is there any downtime for users. Note that Flashback Drop is specifically for the DROP command; you cannot flash back a TRUNCATE command. Along with the table itself, any associated indexes and permissions will also be restored. You cannot flash back a table truncation, only a table drop. Database Control also has an interface to Flashback Drop. From the database home page, take the Availability tab then Perform Recovery in the Manage section. In the drop-down box for Recovery Scope, select TABLES, and then the radio button for Flashback Dropped Tables. `The recycle bin` is a term given to the storage space used by `dropped objects`.

The capacity of flashback drop is limited by the available space in tablespaces

Enabling a table for Flashback Data Archive creates another table (and a few other objects) that will store all previous versions of rows: storing them forever if desired. When DML is committed against the table, a background process called the Flashback `Data Archive process` (the FBDA) will capture the necessary data and save it to the archive. From there, it can be queried with the same syntax used for a regular flashback query - but the flashback can go back years.

The Flashback Data Archive is for long-term storage. Typically, this will be for legal reasons: in many jurisdictions there are requirements for keeping data for years, and then for destroying it (sometimes known as “digital shredding”). A Flashback Data Archive can enable this transparently. When a Flashback Data Archive is created, you specify a time limit (which may be years) and the FBDA will make sure that all data is saved until that time has passed, and will then remove it. The FBDA is responsible for creating the objects in the archive, populating them with rows as necessary and purging data that has passed the expiry date.

There are three data dictionary views that document the Flashback Data Archive configuration:

- `DBA_FLASHBACK_ARCHIVE Describes the configured archives.`
- `DBA_FLASHBACK_ARCHIVE_TS Shows the quotas assigned per archive per tablespace.`
- `DBA_FLASHBACK_ARCHIVE_TABLES Lists the tables for which archiving are enabled.`